

3F4: Data Transmission

Handout 12:

Transmission Control Protocol (TCP)

Congestion Control: TCP-Reno

Ioannis Kontoyiannis

Signal Processing and Communications Lab

Department of Engineering

i.kontoyiannis@eng.cam.ac.uk

Lent Term 2019

1 / 18

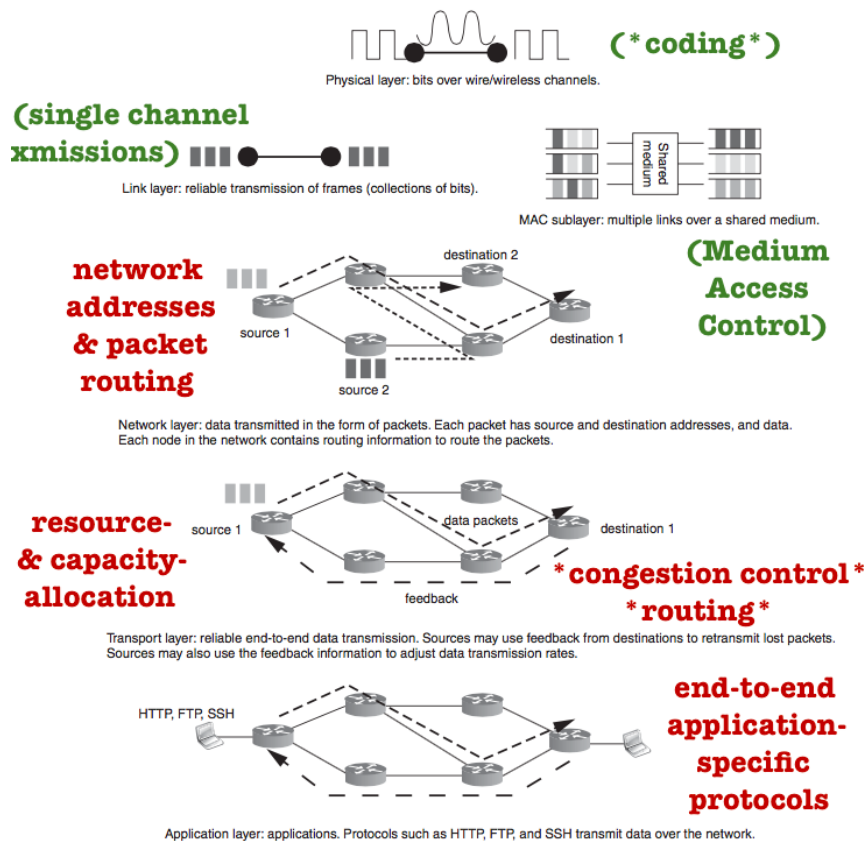


Figure 1.1 Schematic of the layered architecture of a communication network.

Image from: Srikant and Ying "Communication networks: an optimization, control, and stochastic networks perspective." Cambridge University Press, 2013

2 / 18

Transmission Control Protocol (TCP)

- End-to-end transport protocol with many parts
- Used by http (web), smtp (email), telnet, ftp (file transfer), chat, . . .

TCP

1. Connection establishment
2. Connection maintenance
 - Reliability
 - **Congestion control**
 - Flow control
 - Sequencing
3. Connection termination

3 / 18

TCP historical evolution

- 1974** TCP first described (Vint Cerf, Bob Kahn)
- 1975** Three-way handshake (Ray Tomlinson)
- 1981** TCP/IP (Internet Engineering Task Force Standard)
- 1983** BSD Unix 4.2 supports TCP/IP
- 1984** Nagel's algorithm predicts congestion collapse
- 1986** First internet **congestion collapse** observed!
Link Lawrence Berkeley Lab to UC Berkeley
400 yards, 3 hops, throughput 32 Kbps
Dropped to 40 bps: factor of ≈ 1000 drop!
- 1988** Jacobson's algorithms: TCP-Tahoe
slow start, congestion avoidance
- 1990** TCP-Reno
- 1993** TCP-Vegas
- 1996** TCP-NewReno


4 / 18

Congestion

- “too many sources sending too much data too fast for the **network** to handle”
- manifests as **packet loss** (buffer overflow at routers) or **large delay** (long queues at router buffers)
- a very serious issue

Effects of congestion:

- Packet loss
- Retransmission
- Reduced throughput
- Congestion collapse due to
 - Unnecessary retransmissions
 - Undelivered packets
- Congestion may continue after the overload!



5 / 18

Congestion control goals

- High network utilization
- Maximize throughput = total # of bits end-end
- Avoid congestion collapse
- Fair bandwidth sharing
 - Give different sessions equal share
 - Max-min fairness: Maximize the minimum rate session

TCP approach to congestion control

- no explicit feedback from network: closed-loop, end-to-end congestion control
- congestion inferred from end-system observed packet loss and delay
- works well so far: the bandwidth of the Internet has increased by > 5 orders of magnitude

6 / 18

TCP-Reno

One of the most common congestion control algorithms used in the Internet today


It consists of several important elements that are part of almost all common congestion control algorithms

adaptive window flow control
slow start phase
congestion avoidance phase
fast retransmit
fast recovery

7 / 18

TCP-Reno: Preliminaries


- Source sends packets to receiver
- Receiver responds with an acknowledgment **ack** every time a packet is received
- Let T denote the **RTT**: $T = \text{round-trip time}$ between packet being sent and **ack** received
- Make repeated measurements of time from packet transmission until **ack** receipt
- Compute empirical mean μ and standard deviation σ
- Set $T = \mu$
- Set **timeout** $= \mu + (\text{a few})\sigma$



8 / 18

TCP-Reno: Window flow control and **ack**'s

- Source maintains a **window** containing a varying number of W packets
- Every time an **ack** is received, window shifts by 1 position
- An **ack** contains the index of the next packet the receiver expects. E.g., if it's received packets 1,2, upon receiving packet 3 the next **ack** asks for packet 4
- If then packet 5 is received next **ack** asks for packet 4 again
- This is called a 'duplicate acknowledgment' or **dupack**



9 / 18

TCP-Reno: Initialization

Initialize:

- Set window size $W = 1$
- Estimate RTT T
- Select slow start threshold value **ssthresh** for W

Basic TCP-Reno operation:

- Send all W packets currently in the window
- Wait T seconds for the corresponding **ack**'s
- Update $W \mapsto W'$ and **ssthresh**
- Send next W' packets in the new window

Transmission rate:

- At time t , the rate $R(t) = \frac{W(t)}{T}$ packets/second

10 / 18


TCP-Reno: Slow start phase

Idea: Increase W by 1
with each **ack**

While all is going well:

- Send $W = 1$ packet
receive **ack**, set $W = 2$
- Send $W = 2$ packets
receive 2 **ack**'s, set $W = 4$
- Send $W = 4$ packets
receive 4 **ack**'s, set $W = 8$
- Continue until
 $W > \text{ssthresh}$

When $W > \text{ssthresh}$
enter congestion avoidance phase



11 / 18

TCP-Reno: Congestion avoidance phase

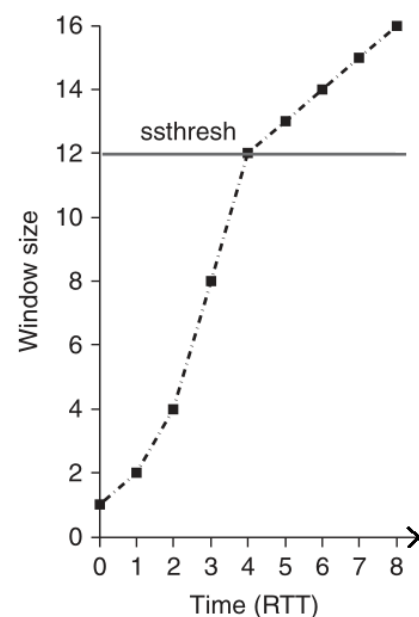
Idea: Increase W
by $1/W$ with each **ack**

While all is going well:

- Send W packets
receive W **ack**'s
set $W \mapsto W + 1$
- Continue until
something goes wrong


What can go wrong?

- **Delay** = three **dupack**'s
- **Loss** = **timeout**



12 / 18

TCP-Reno: Three **dupack**'s \Rightarrow 'Fast Retransmit'



Transmitter detects **delay**:


- If sender receives 3 **dupack**'s for the same data it assumes that packet was irreparably delayed
- Resend missing packet before **timeout** (hence 'fast retransmit')
- Change the value of **ssthresh** to $W/2$
- Set the window size $W \mapsto W/2$
- Continue as before

13 / 18

TCP-Reno: **timeout** \Rightarrow 'Fast Recovery'

Transmitter detects **loss**:

- If sender receives no **ack**'s for any of the packets sent from the window by **timeout**, it assumes they are permanently lost
- Change the value of **ssthresh** to $W/2$
- Set window size $W = 1$
- Return to slow start phase: Start by resending lost packets



14 / 18

TCP-Reno: Observations

Revisit five TCP-Reno elements:

- Variable **window size** regulates transmission speed
- **Slow start**: Multiplicative (i.e., exponential) rate increase until $ssthresh/T$ packets/second
- **Congestion avoidance**: Additive (i.e., linear) rate increase until delay/loss detected
- **Fast retransmit** when delay detected
- **Fast recovery** when **timeout**/loss detected

Observe:

- Loss (**timeout**) is treated as a more serious issue than delay (3 **dupack**'s)
- In both cases, multiplicative rate decrease
- In the long run, TCP-Reno spends almost all of the time in the congestion avoidance phase
- Additive Increase Multiplicative Decrease (AIMD) adaptation

15 / 18

TCP-Reno: Equilibrium rate

What is the algorithm's
long-term average transmission rate $R(t)$?

- Recall rate $R(t) = W(t)/T$
- Assume in CA phase (ignore SS)
- Consider a continuous-time approximate model
- Let $q(t) =$ loss rate at time t

In the long run, how does $R(t)$ depend on $q(t)$?

Strategy:

- Compute the rate of change $W'(t)$ of the window size
- Assume that, at equilibrium, $R'(t) = W'(t)/T = 0$
- Solve $R'(t) = 0$

16 / 18

TCP-Reno: Continuous-time model

- rate at which **ack**'s are received is $R(t - T)(1 - q(t))$
- each **ack** increases $W(t)$ by $1/W(t)$
- packet loss rate is $R(t - T)q(t)$
- each loss event decreases $W(t)$ by a factor $\beta = 1/2$

Therefore:

$$W'(t) = \frac{R(t - T)(1 - q(t))}{W(t)} - \beta R(t - T)q(t)W(t)$$

i.e.,
$$R'(t) = \frac{R(t - T)(1 - q(t))}{T^2 R(t)} - \beta R(t - T)q(t)R(t)$$

and solving $R'(t) = 0$:

TCP square-root law


$$R(t) = \frac{1}{T\sqrt{\beta}} \sqrt{\frac{1 - q(t)}{q(t)}}$$

17 / 18

TCP-Reno in equilibrium: Conclusions

$$R(t) = \frac{1}{T\sqrt{\beta}} \sqrt{\frac{1 - q(t)}{q(t)}}$$

- Rate $R(t) \propto 1/\text{RTT}$
- When the loss probability $q(t)$ is small the rate $R(t) \propto 1/\sqrt{q(t)}$
- Results extensively verified in empirical observations
E.g., 3 groups of 10 sources each, $3\text{ms} \leq \text{RTT} \leq 7\text{ms}$,
link capacity = 64Mbps



18 / 18